
CS2030 Lecture 9

Java Streams

Henry Chia (hchia@comp.nus.edu.sg)

Semester 1 2023 / 2024

Lecture Outline and Learning Outcomes

- Know how to create **stream** pipelines for *internal* iteration
 - Know the difference between primitive and generic streams
- Understand **lazy evaluation** in source/intermediate operations, and **eager evaluation** for terminal operations
- Appreciate how lazy evaluation supports **infinite stream**
- Able to implement a basic lazy context by encapsulating a `Supplier` functional interface for *delayed data*
- Appreciate that streams should be *inherently parallelizable*
- Know how to write correct streams that are non-interfering and stateless with no side effects

External Iteration

- An external iteration is defined *imperatively*
 - e.g. sum of all integers in the closed interval $[1, 10]$

```
jshell> int sum = 0  
sum ==> 0
```

```
jshell> for (int x = 1; x <= 10; x = x + 1) {  
    ...>     sum = sum + x;  
    ...> }
```

```
jshell> sum  
sum ==> 55
```

- Errors could be introduced when
 - sum is initialized wrongly before the loop
 - looping variable x is initialized wrongly
 - loop condition is wrong
 - increment of x is wrong
 - aggregation of sum is wrong

Internal Iteration: Stream

- Internal iteration is defined *declaratively*
 - e.g. using a primitive integer stream

```
jsell> int sum = IntStream.rangeClosed(1, 10).  
...> sum()  
sum ==> 55
```
- Literal meaning “loop through values 1 to 10, and sum them”
- No need to specify how to iterate through elements or use any *mutable* variables — no variable state, no surprises! 😊
- A **stream** is a sequence of elements on which tasks are performed; stream elements move through a sequence of tasks in the stream pipeline
- Result is obtained at the end of stream processing

Stream Pipeline

- A stream pipeline comprises
 - a **data source** (e.g. `IntStream::rangeClosed`) to start the stream
 - some **intermediate operations** (e.g. `IntStream::map`) that specify tasks to perform on a stream's elements

```
jshell> IntStream stream = IntStream.rangeClosed(1, 10).
...> map(x -> x * 2)
stream ==> java.util.stream.IntPipeline$Head@12edcd21
```
 - a **terminal operation** (e.g. `IntStream::sum`) that *reduces* the stream elements into a single value

```
jshell> stream.sum()
$.. ==> 110
```
- Each source/intermediate operation returns a new stream of processing steps specified up to that point in the pipeline
- Stream elements within a stream *can only be consumed once*

Exercise: Primality Test

- Given the following external iteration:

```
boolean isPrime(int n) {  
    for (x = 2; x < n; x++) { // x <= (int) Math.sqrt(n)  
        if (n % x == 0) {  
            return false;  
        }  
    }  
    return true;  
}
```

- Complete the following internal iteration:

```
boolean isPrime(int n) {  
    return n > 1 && IntStream...  
  
}
```

Reducing a Stream to a Value

- Iterate through `IntStream` elements and reduce to an `int`
`int reduce(int identity, IntBinaryOperator op)`
- `IntBinaryOperator` with single abstract method:
`int applyAsInt(int left, int right)`

```
jshell> IntStream.rangeClosed(1, 10).  
...> reduce(0, (x,y) -> x + y)  
$.. ==> 55  
  
jshell> IntStream.rangeClosed(1, 10).  
...> reduce(1, (x,y) -> x * y)  
$.. ==> 3628800
```
- Alternative one argument reduce that returns `OptionalInt`
`OptionalInt reduce(IntBinaryOperator op)`

```
jshell> IntStream.range(1, 10).reduce((x, y) -> x * y)  
$.. ==> OptionalInt[362880]  
  
jshell> IntStream.range(1, 1).reduce((x, y) -> x * y)  
$.. ==> OptionalInt.empty
```

flatMap Method in Stream

- How about nested loops?

```
for (x = 1; x <= 3; x++)  
    for (y = x; y <= 3; y++)  
        System.out.println((x * y) + " "); // output is 1 2 3 4 6 9
```

- map tries to map each stream element into one other stream

```
jshell> IntStream.rangeClosed(1, 3).  
...> map(x -> IntStream.rangeClosed(x,3).map(y -> x * y))  
| Error:  
| incompatible types: bad return type in lambda expression  
|     java.util.stream.IntStream cannot be converted to int  
| map(x -> IntStream.rangeClosed(x,3).map(y -> x * y))  
|     ^-----^
```

- flatMap transforms each stream element into a stream of other elements (either zero or more) by taking in a function that produces another stream, and then *flattens* it

```
jshell> IntStream.rangeClosed(1, 3).  
...> flatMap(x -> IntStream.rangeClosed(x,3).map(y -> x * y)).  
...> forEach(x -> System.out.print(x + " "))  
1 2 3 4 6 9
```


Generic Stream<T>

- Stream<T> is a stream over reference-typed objects with data sources: of, iterate and generate

```
jshell> int sum = Stream.<Integer>iterate(1, x -> x <= 10, x -> x + 1).  
    ...> reduce(0, (x,y) -> x + y) // note: reduce(T, BinaryOperator<T>)  
sum ==> 55
```

- boxed() wraps stream elements in its wrapper type

```
jshell> Stream<Integer> stream = IntStream.rangeClosed(1, 10).boxed()  
stream ==> java.util.stream.IntPipeline$1@5010be6  
  
jshell> List<Integer> list = stream.toList()  
list ==> [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

- mapToObj converts from primitive to generic stream

```
jshell> IntStream.rangeClosed(1, 10).mapToObj(x -> "<" + x + ">").  
    ...> toList()  
$.. ==> [<1>, <2>, <3>, <4>, <5>, <6>, <7>, <8>, <9>, <10>]
```

- Stream::toList() converts generic stream to generic list
- List::stream() converts generic list to generic stream

Lazy Evaluation in Streams

- Source/intermediate operations use **lazy evaluation**
 - does not perform any operations on stream's elements until a terminal operation is called
- Terminal operations use **eager evaluation**
 - performs the requested operation as soon as it is called

```
jshell> Stream.<Integer>iterate(1, x -> x + 1).
...> limit(5).
...> peek(x -> System.out.println("limit: " + x)).
...> filter(x -> x % 2 == 0).
...> peek(x -> System.out.println("filter: " + x)).
...> map(x -> x * 2).
...> peek(x -> System.out.println("map: " + x)).
...> reduce(0, (x, y) -> {
...>     System.out.println("reduce: " + x + " + " + y);
...>     return x + y;
...> })
```

limit: 1
limit: 2
filter: 2
map: 4
reduce: 0 + 4
limit: 3
limit: 4
filter: 4
map: 8
reduce: 4 + 8
limit: 5
\$.. ==> 12

Infinite Stream

- Lazy evaluation allows us to work with infinite streams that represent an infinite number of elements
 - `Stream<T>::generate(Supplier<T> supplier)` produces an infinite sequence of values generated by `supplier`
 - `Stream<T>::iterate(T seed, UnaryOperator<T> next)` produces an infinite sequence by repeatedly applying the function `next` starting with the `seed` value
- Intermediate operations, e.g. `limit`, can be used to restrict the total number of elements in the stream

```
jshell> Stream.<Integer>iterate(1, x -> x + 1).  
...> filter(x -> x % 2 == 1).  
...> limit(20). // find first 20 odd numbers  
...> forEach(x -> System.out.print(x + " "))  
1 3 5 7 9 11 13 15 17 19 21 23 25 27 29 31 33 35 37 39
```

Lazy Class

- To understand how lazy evaluation works, define a Lazy class

```
import java.util.function.Supplier;

class Lazy<T> {
    private final Supplier<T> supplier;

    private Lazy(Supplier<T> supplier) {
        this.supplier = supplier;
    }

    static <T> Lazy<T> of(Supplier<T> supplier) {
        return new Lazy<T>(supplier);
    }

    static <T> Lazy<T> of(T t) {
        return new Lazy<T>(() -> t);
    }

    public T get() {
        return supplier.get();
    }
}
```

```
jshell> int foo() {
...>     System.out.println("foo");
...>     return -1;
...> }
| created method foo()
jshell> Lazy<Integer> lazy = Lazy.of(foo())
foo
$.. ==> Lazy@ae45eb6
jshell> lazy.get()
$.. ==> -1
jshell> lazy = Lazy.<Integer>of(() -> foo())
$.. ==> Lazy@6f7fd0e6
jshell> lazy.get()
foo
$.. ==> -1
```

- `Lazy.of(foo())` evaluates `foo` method *eagerly*
- `Lazy.of(() -> foo())` evaluates `foo` *lazily*, i.e. only when `get()` is invoked sometime later

Mapping a Lazy Value

- Define map that returns a new Lazy

```
<R> Lazy<R> map(Function<? super T, ? extends R> mapper) {  
    Supplier<R> supplier = () -> mapper.apply(this.get());  
    return Lazy.<R>of(supplier);  
}
```

```
jshell> Lazy<Integer> i = Lazy.<String>of(() -> "abc").  
...> map(x -> { System.out.println("map1"); return x.length(); }).  
...> map(x -> { System.out.println("map2"); return x * 2; })  
i ==> Lazy@51565ec2 // map is not evaluated until a get()  
  
jshell> i.get() // map is lazily evaluated :)  
map1  
map2  
$.. ==> 6
```

- What is wrong with the following implementation of map?

```
<R> Lazy<R> map(Function<? super T, ? extends R> mapper) {  
    R r = mapper.apply(this.get());  
    return Lazy.<R>of(() -> r);  
}
```

Inherently Parallelizable Stream

- A stream pipeline should be *inherently parallelizable*
 - intermediate operations can operate on elements in parallel
 - reduction uses a divide-and-conquer strategy
- `parallel()` operation switches the stream pipeline to parallel
 - invoke anywhere between the data source and terminal
 - `sequential()` switches off parallel operation
- `reduce` method uses an *associative accumulation function*, e.g.

```
T reduce(T identity, BinaryOperator<T> acc) // BiFunction<T,T,T>
```

```
jshell> DoubleStream.of(1.0, 2.0, 3.0, 4.0).parallel().  
...> reduce(1.0, (x, y) -> x * y) // multiply is associative  
$.. ==> 24.0
```

```
jshell> DoubleStream.of(1.0, 2.0, 3.0, 4.0).parallel().  
reduce(24.0, (x, y) -> x / y) // divide is not associative  
$.. ==> 1.5 // should be 1.0?
```

Example: Parallelism in Streams

□ Parallelizing the search for primes

```
jshell> Runtime.getRuntime().availableProcessors()
$.. ==> 8

jshell> ForkJoinPool.commonPool().getParallelism()
$.. ==> 7

jshell> import java.time.*
jshell> long numOfPrimes(int from, int to) {
...>     Instant start = Instant.now(); // start timing
...>     long howMany = IntStream.rangeClosed(from, to)
...>         .parallel()
...>         .filter(x -> isPrime(x))
...>         .count();
...>     Instant stop = Instant.now(); // end timing
...>     System.out.println("Duration: " +
...>         Duration.between(start, stop).toMillis() + "ms");
...>     return howMany;
...> }
jshell> numOfPrimes(2_000_000, 3_000_000)
Duration: 239ms
$.. ==> 67883
```

□ Avoid parallelizing trivial tasks, e.g. parallelizing isPrime

- creates more work in terms of parallelizing overhead
- worthwhile only if the parallel task is complex enough

Correctness of Streams

□ To ensure correct execution, stream operations

– must not interfere with stream data

```
jshell> List<String> list = new ArrayList<String>(  
...> List.of("abc", "def", "xyz"))  
list ==> [abc, def, xyz]  
  
jshell> list.stream().peek(str -> {  
...> if (str.equals("xyz")) { list.add("pqr"); }  
...> }).forEach(x -> {})  
| Exception java.util.ConcurrentModificationException  
| ...
```

– *preferably* stateless (map vs distinct) with no side effects

```
jshell> List<Integer> list = List.of(1, 3, 5, 7, 9, 11, 13, 15, 17, 19)  
list ==> [1, 3, 5, 7, 9, 11, 13, 15, 17, 19]  
  
jshell> List<Integer> result = new ArrayList<Integer>()  
result ==> []  
  
jshell> list.stream().parallel(). // or list.parallelStream().  
...> filter(x -> isPrime(x)).  
...> forEach(x -> result.add(x)) // what is result?
```